


Class 7 Machine Learning 1

Kyle Alvarez

K-means clustering

First we will test how this method works in R with some made up data.

```
x <- rnorm(10000)  
hist(x)
```




Let's make some numbers centered on -3

```
rev(c("a", "b", "c"))
```

```
[1] "c" "b" "a"
```

```
tmp <- c(rnorm(30, -3), rnorm(30, +3))  
  
x <- cbind(x= tmp, y= rev(tmp))  
plot(x)
```



Now, let's see how `kmeans()` works with this data...

```
km <- kmeans(x, centers = 2, nstart=20)  
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

| | x | y |
|---|-----------|-----------|
| 1 | 3.027637 | -2.988773 |
| 2 | -2.988773 | 3.027637 |

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
```

```
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 56.69148 56.69148
(between_SS / total_SS = 90.5 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"       "withinss"    "tot.withinss"
[6] "betweenss"    "size"        "iter"       "ifault"
```

```
km$centers
```

```
      x      y
1  3.027637 -2.988773
2 -2.988773  3.027637
```

Q. How many points are in each cluster?

```
km$size # cluster size
```

```
[1] 30 30
```

Q. What 'component' of your result object details - cluster size? - cluster assignment/membership? - cluster center?

```
km$cluster #cluster assignment/membership
```


```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

```
km$centers #cluster centers
```

```
      x      y
1  3.027637 -2.988773
2 -2.988773  3.027637
```

Q. Plot x colored by the kmeans cluster assignment and as cluster centers as blue points

```
plot(x, col = km$cluster)
points(km$centers, col = "blue", pch=15, cex = 2)
```



```
#pch changes the shape
#cex makes it bigger or smaller
```

Hierarchical Clustering

The `hclust()` function in R performs hierarchical clustering

The `hclust()` function requires an input distance matrix, which I can get from the `dist()` function.

```
hc <- hclust(dist(x))
hc
```

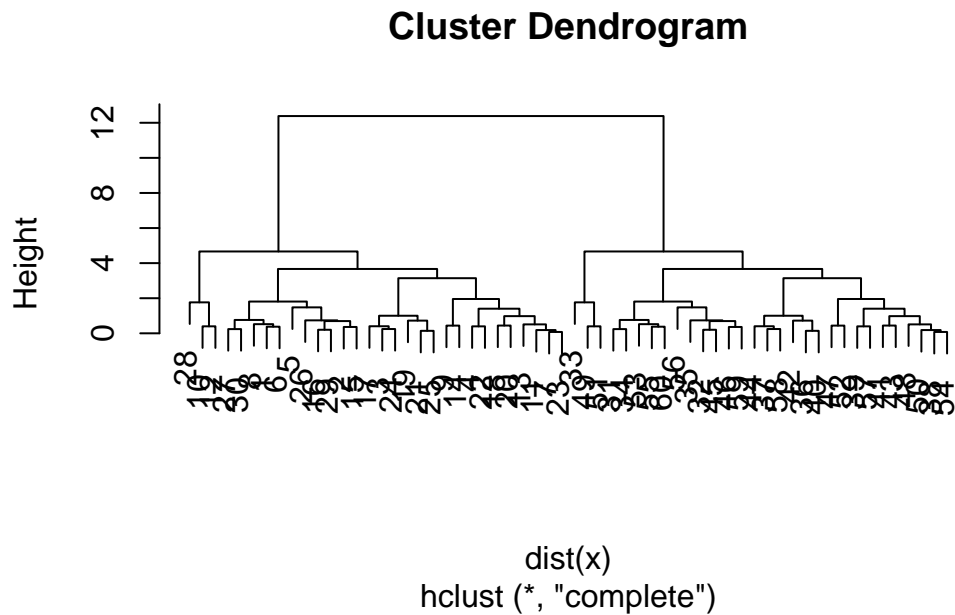
Call:

```
hclust(d = dist(x))
```

```
Cluster method   : complete
Distance        : euclidean
Number of objects: 60
```

There is a `plot()` method for `hclust` objects...

```
plot(hc)
```



Now to get my cluster membership vector I need to “cut” the tree to yield separate “branches” with the “leaves” on each branch being our cluster. To do this we use the `cutree()` function.

```
cutree(hc, h=8)
```


[illegible]

Use `cutree()` with a `k=2`.

```
grps <- cutree(hc, k=2)
```

A plot of our data colored by our hclust grps.

```
plot(x, col = grps)
```



Principal Component Analysis (PCA)

First, import data of “UK_foods”

```
url <- "https://tinyurl.com/UK-foods"  
foods <- read.csv(url, row.names = 1)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

Can use the `dim()` function which returns the number of rows and columns of a data frame

```
dim(foods)
```

```
[1] 17  4
```

Now we should check our data using the `view()` function

```
View(foods)
```

Clean up our data frame as we were only expecting 4 columns (one for each country in the UK).


```
## Not good
##rownames(foods) <- foods[,1]
##foods <- foods[, -1]
##head(foods)
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

The approach in which we manually delete the column name isn’t as good because everytime that chunk of code runs, it will delete each column as it runs. So using the `row.names=1` argument in the `read.csv()` function is preferable as it doesn’t do this.


Can use a bar plot to plot the data above

```
barplot(as.matrix(foods), beside=T, col=rainbow(nrow(foods)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

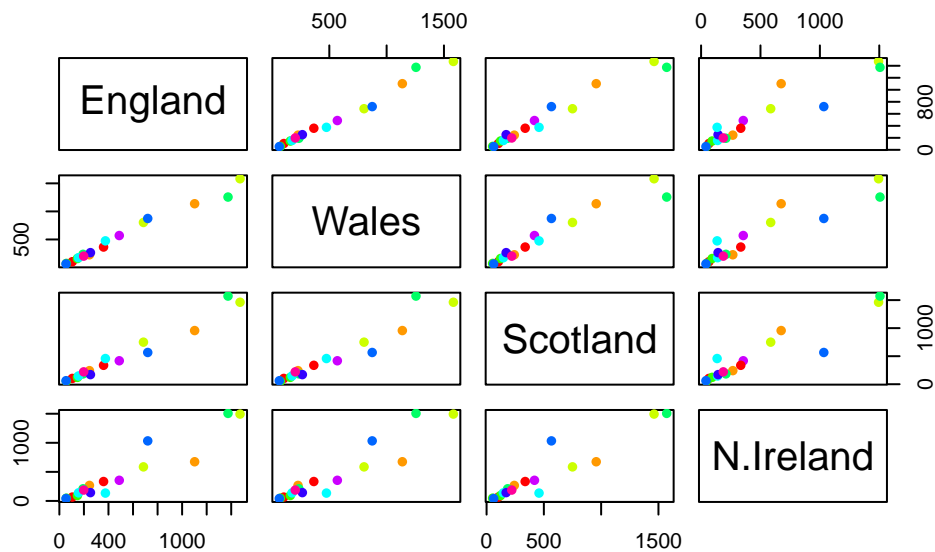
```
barplot(as.matrix(foods), beside=F, col=rainbow(nrow(foods)))
```



The optional argument changes the barplot to stack the different foods on top eachother for each country.

Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(foods, col=rainbow(10), pch=16)
```

If a point lies on the diagonal for a given plot, that means they are positively correlated, which also means they are essentially the same amount. The following code above, basically makes a pair plot for one country vs another country, for all the countries.

While this is kind of useful, it takes work to dig into the details here to find out what is different in these countries.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

The main different between N. Ireland and other countries of the UK in terms of this data-set is that there is a blue dot that is higher in N. Ireland compared to other countries.

PCA to the Rescue

Principal Component Analysis (PCA for short) can be a big help in these cases where we have lot's of things that are being measured in a dataset.

The main PCA function in base R is called `prcomp()`.

Note: The `prcomp()` function wants as input the transpose of our food matrix/table/data.frame

```
pca <- prcomp( t(foods) )
summary(pca)
```

Importance of components:

| | PC1 | PC2 | PC3 | PC4 |
|------------------------|----------|----------|----------|-----------|
| Standard deviation | 324.1502 | 212.7478 | 73.87622 | 4.189e-14 |
| Proportion of Variance | 0.6744 | 0.2905 | 0.03503 | 0.000e+00 |
| Cumulative Proportion | 0.6744 | 0.9650 | 1.00000 | 1.000e+00 |

The above result shows that PCA captures 67% of the total variance in the original data in one PC and 96.5% in two PCs.

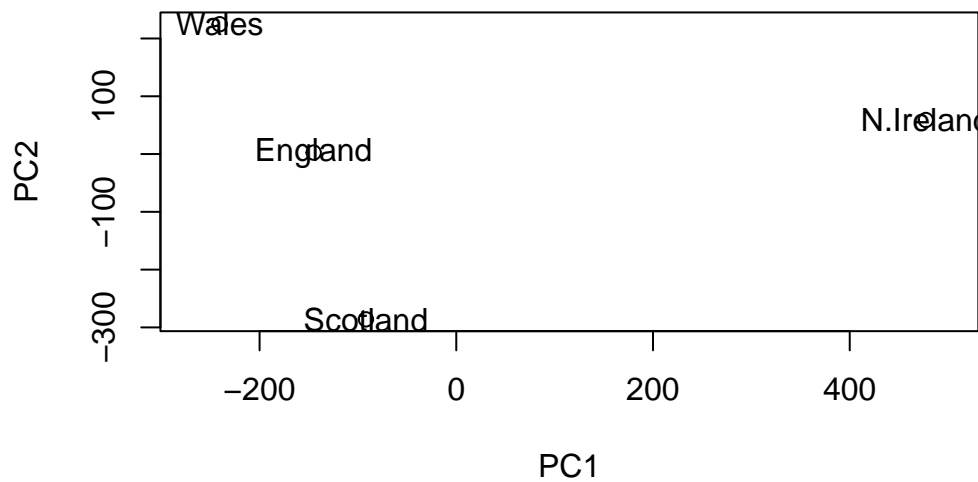
```
head(pca$x)
```

| | PC1 | PC2 | PC3 | PC4 |
|-----------|------------|-------------|-------------|---------------|
| England | -144.99315 | 2.532999 | -105.768945 | 2.842865e-14 |
| Wales | -240.52915 | 224.646925 | 56.475555 | 7.804382e-13 |
| Scotland | -91.86934 | -286.081786 | 44.415495 | -9.614462e-13 |
| N.Ireland | 477.39164 | 58.901862 | 4.877895 | 1.448078e-13 |

Let's plot our main results.


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(foods))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
#Color our plot
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500), col=c("orange", "red", "blue", "green"))
text(pca$x[,1], pca$x[,2], colnames(foods), col=c("orange", "red", "blue", "green"))
```



Calculate how much standard deviation to calculate how much variation in the original data each PC accounts for.

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```


```
[1] 67 29 4 0
```

```
## or the second row here...
z <- summary(pca)
z$importance
```

| | PC1 | PC2 | PC3 | PC4 |
|------------------------|-----------|-----------|----------|--------------|
| Standard deviation | 324.15019 | 212.74780 | 73.87622 | 4.188568e-14 |
| Proportion of Variance | 0.67444 | 0.29052 | 0.03503 | 0.000000e+00 |
| Cumulative Proportion | 0.67444 | 0.96497 | 1.00000 | 1.000000e+00 |

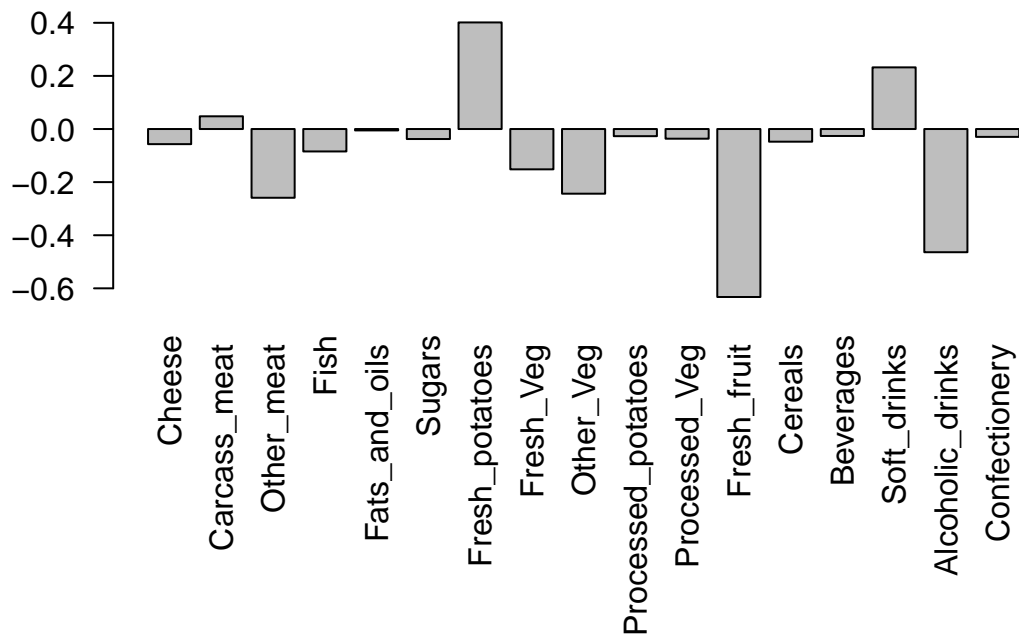
We can summarize the plot of variances (eigenvalues) with respect to the principal component number (eigenvector number)

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



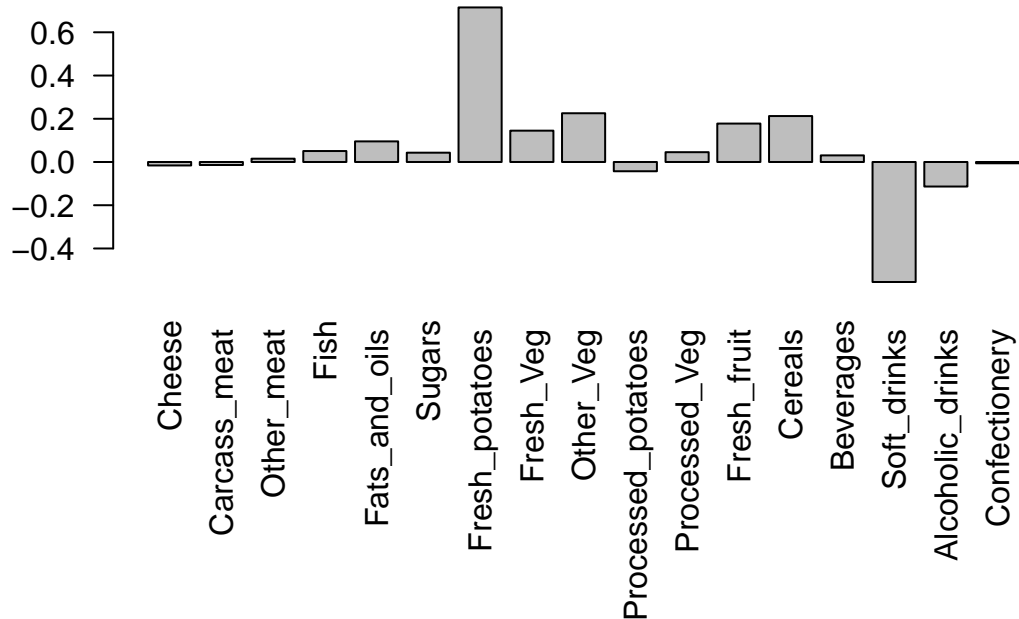
We can consider the influence of each of the original variables upon the principal components (loading scores). Can be summarized with a call to `biplot()`

```
## Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?


```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



The two food groups that are featured prominently are fresh potatoes and soft drinks. PC2 mainly tells us the second most variation.

Can use a biplot to visualize PCA (main way):

```
## The inbuilt biplot() can be useful for small datasets
biplot(pca)
```



PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

| | wt1 | wt2 | wt3 | wt4 | wt5 | ko1 | ko2 | ko3 | ko4 | ko5 |
|-------|------|-----|------|------|-----|-----|-----|-----|-----|-----|
| gene1 | 439 | 458 | 408 | 429 | 420 | 90 | 88 | 86 | 90 | 93 |
| gene2 | 219 | 200 | 204 | 210 | 187 | 427 | 423 | 434 | 433 | 426 |
| gene3 | 1006 | 989 | 1030 | 1017 | 973 | 252 | 237 | 238 | 226 | 210 |
| gene4 | 783 | 792 | 829 | 856 | 760 | 849 | 856 | 835 | 885 | 894 |
| gene5 | 181 | 249 | 204 | 244 | 225 | 277 | 305 | 272 | 270 | 279 |
| gene6 | 460 | 502 | 491 | 491 | 493 | 612 | 594 | 577 | 618 | 638 |

Q10: How many genes and samples are in this data set?

```
dim(rna.data)
```


```
[1] 100 10
```


There is 100 genes and 10 samples.

Can generate our data using PCA

```
## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)

## Simple un polished plot of pc1 and pc2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```



Summary of how much variation in the original data PC accounts for:

```
summary(pca)
```

Importance of components:


| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 | PC7 |
|------------------------|--------|--------|---------|---------|---------|---------|---------|
| Standard deviation | 9.6237 | 1.5198 | 1.05787 | 1.05203 | 0.88062 | 0.82545 | 0.80111 |
| Proportion of Variance | 0.9262 | 0.0231 | 0.01119 | 0.01107 | 0.00775 | 0.00681 | 0.00642 |
| Cumulative Proportion | 0.9262 | 0.9493 | 0.96045 | 0.97152 | 0.97928 | 0.98609 | 0.99251 |

| | PC8 | PC9 | PC10 |
|------------------------|---------|---------|-----------|
| Standard deviation | 0.62065 | 0.60342 | 3.348e-15 |
| Proportion of Variance | 0.00385 | 0.00364 | 0.000e+00 |

Cumulative Proportion 0.99636 1.00000 1.000e+00

Can use a barplot summary of the Proportion of Variance for each PC

```
plot(pca, main="Quick scree plot")
```



Can use the square of `pca$sdev` (standard deviation), to calculate how much variation in the original data each PC accounts for


```
## Variance captured per PC
pca.var <- pca$sdev^2

## Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
[1] 92.6 2.3 1.1 1.1 0.8 0.7 0.6 0.4 0.4 0.0
```

Then using these calculations, we can generate a scree-plot:

```
barplot(pca.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```




Can color our PCA plot to make it more aesthetic:

```
## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
      xlab=paste0("PC1 (", pca.var.per[1], "%)"),
      ylab=paste0("PC2 (", pca.var.per[2], "%)"))

text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```




Using ggplot2

Can use ggplot to plot our PCA results:

```
library(ggplot2)

df <- as.data.frame(pca$x)


# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```



Label and color for wild-type and knock-out samples.

```
# Add a 'wt' and 'ko' "condition" column
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
p
```




A more polish looked at our graph:

```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="Class example data") +
theme_bw()
```

PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



Class example data

Top 10 measurements that contribute to pc1 in either direction (+ or -):

```
loading_scores <- pca$rotation[,1]

## Find the top 10 measurements (genes) that contribute
## most to PC1 in either direction (+ or -)
gene_scores <- abs(loading_scores)
gene_score_ranked <- sort(gene_scores, decreasing=TRUE)

## show the names of the top 10 genes
top_10_genes <- names(gene_score_ranked[1:10])
top_10_genes
```

```
[1] "gene100" "gene66" "gene45" "gene68" "gene98" "gene60" "gene21"
[8] "gene56" "gene10" "gene90"
```