

Class 12: RNASeq Analysis

Kyle Alvarez

In today's class we will work with published RNA-seq experiment where airway smooth muscle cells were treated with dexamethasone, a synthetic glucocorticoid steroid with anti-inflammatory effects (Himes et al. 2014).

Data import

We will use good old `read.csv()` to read the two things we need for this analysis:

- counts data
- col data (meta data)

)

```
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

Q1. How many genes are in this dataset? There are 38694 genes in this dataset.

How many transcripts do I have?

```
nrow(counts)
```

```
[1] 38694
```

Lets have a look at the metadata...

```
metadata
```

```

      id      dex celltype     geo_id
1 SRR1039508 control    N61311 GSM1275862
2 SRR1039509 treated    N61311 GSM1275863
3 SRR1039512 control    N052611 GSM1275866
4 SRR1039513 treated    N052611 GSM1275867
5 SRR1039516 control    N080611 GSM1275870
6 SRR1039517 treated    N080611 GSM1275871
7 SRR1039520 control    N061011 GSM1275874
8 SRR1039521 treated    N061011 GSM1275875

```

And the counts data

```
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG000000000003	723	486	904	445	1170
ENSG000000000005	0	0	0	0	0
ENSG000000000419	467	523	616	371	582
ENSG000000000457	347	258	364	237	318
ENSG000000000460	96	81	73	66	118
ENSG000000000938	0	0	1	0	2
	SRR1039517	SRR1039520	SRR1039521		
ENSG000000000003	1097	806	604		
ENSG000000000005	0	0	0		
ENSG000000000419	781	417	509		
ENSG000000000457	447	330	324		
ENSG000000000460	94	102	74		
ENSG000000000938	0	0	0		

First we should check the correspondence of the metadata and count data

```
metadata$id
```

```
[1] "SRR1039508" "SRR1039509" "SRR1039512" "SRR1039513" "SRR1039516"
[6] "SRR1039517" "SRR1039520" "SRR1039521"
```

```
colnames(counts)
```

```
[1] "SRR1039508" "SRR1039509" "SRR1039512" "SRR1039513" "SRR1039516"
[6] "SRR1039517" "SRR1039520" "SRR1039521"
```

To check that these are all in the same order we can use `==` to test of equality

```
all(metadata$id == colnames(counts))
```

```
[1] TRUE
```

Analysis via comparison of CONTROL vs TREATED

The “treated” have the dex drug and the “control” do not. First I need to be able to extract just the “control” columns in the `counts` data.

```
control inds <- metadata$dex == "control"  
control <- metadata[control inds,]  
control$id
```

```
[1] "SRR1039508" "SRR1039512" "SRR1039516" "SRR1039520"
```

Q2. How many ‘control’ cell lines do we have? We have 4 control cell lines.

Q3. How would you make the above code in either approach more robust? In either to make the above code more robust, you can remove the intermediate step of saving things into variables and essentially call the necessary functions on top of each other. for example you can do `metadata[metadata$dex == “control”, id]` and save this into a variable and then call `rowSums` on the saved variable..

Now I can use this to access just the “control” columns of my `counts` data...

```
control.counts <- counts[,control$id]  
head(control.counts)
```

	SRR1039508	SRR1039512	SRR1039516	SRR1039520
ENSG000000000003	723	904	1170	806
ENSG000000000005	0	0	0	0
ENSG00000000419	467	616	582	417
ENSG00000000457	347	364	318	330
ENSG00000000460	96	73	118	102
ENSG00000000938	0	1	2	0

Find the mean count value for each transcript/gene by binding the `rowMeans()`.

```
control.mean <- rowMeans(control.counts)
head(control.mean)
```

```
ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
         900.75          0.00        520.50        339.75        97.25
ENSG000000000938
         0.75
```

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called treated.mean).

And now find a mean value for all the “treated” columns in the same way

```
treated.id <- metadata[metadata$dex == "treated", "id"]
treated.mean <- rowMeans(counts[,treated.id])
```

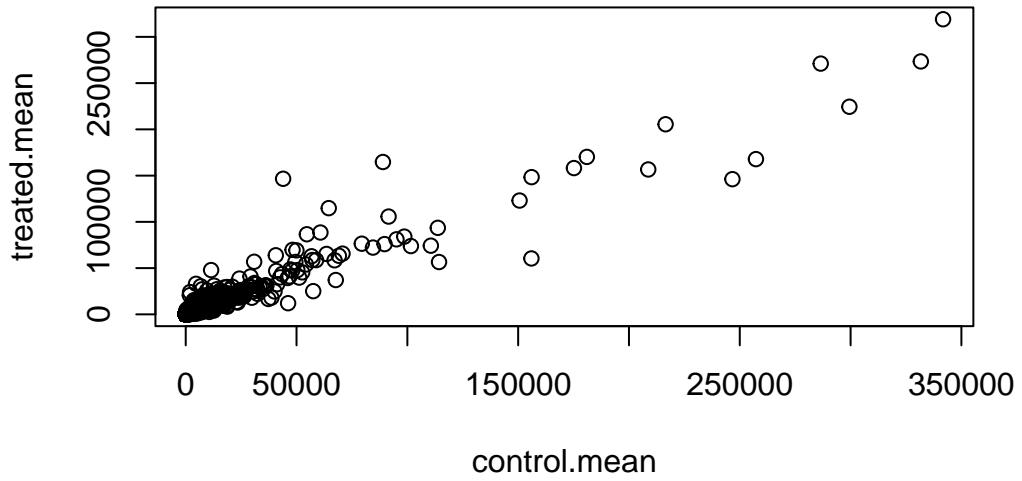
Now I have control.mean and treated.mean. Lets put them together for safe keeping and ease of use later.

```
meancounts <- data.frame(control.mean, treated.mean)
head(meancounts)
```

	control.mean	treated.mean
ENSG000000000003	900.75	658.00
ENSG000000000005	0.00	0.00
ENSG000000000419	520.50	546.00
ENSG000000000457	339.75	316.50
ENSG000000000460	97.25	78.75
ENSG000000000938	0.75	0.00

Q5 (a). Create a scatter plot showing the mean of the treated samples against the mean of the control samples. Your plot should look something like the following. Q5 (b). You could also use the ggplot2 package to make this figure producing the plot below. What geom_?() function would you use for this plot? Use geom_point() to make the plot in ggplot. Let’s do a quick plot to see how our data looks

```
plot(meancounts)
```

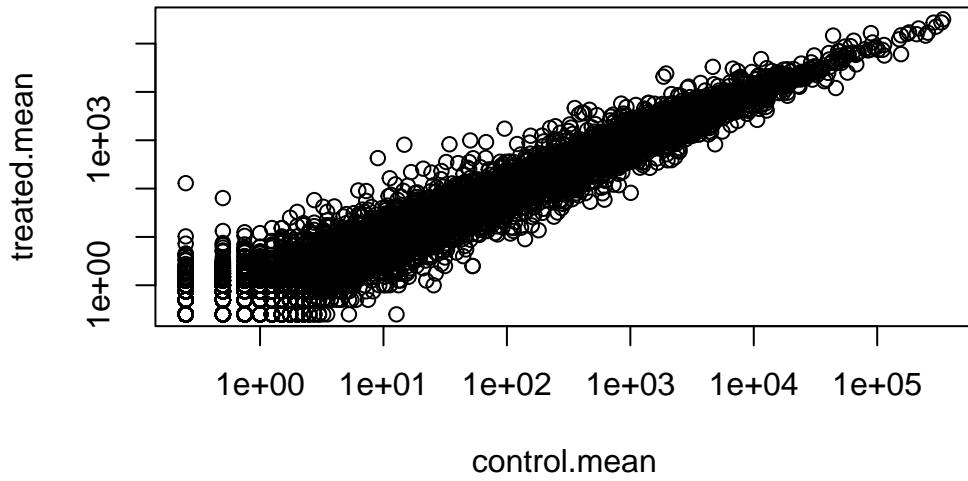


Q6. Try plotting both axes on a log scale. What is the argument to plot() that allows you to do this? Using the log argument in the plot function. This is a very heavily skewed and over a wide range - calling out for a log transform!

```
plot(meancounts, log="xy")
```

Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted from logarithmic plot

Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted from logarithmic plot



We like working with log transformed data as it can help make things more straightforward to interpret.

If we have no change:

```
log2(20/20)
```

```
[1] 0
```

What about if we had a doubling:

```
log2(40/20)
```

```
[1] 1
```

Half as much

```
log2(10/20)
```

```
[1] -1
```

```
log2(80/20)
```

```
[1] 2
```

We like working with log2 fold-change values. Let's calculate them for our data.

```
meancounts$log2fc <- log2(meancounts$treated.mean/meancounts$control.mean)
head(meancounts)
```

	control.mean	treated.mean	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000005	0.00	0.00	NaN
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000938	0.75	0.00	-Inf

Q7. What is the purpose of the arr.ind argument in the which() function call above? Why would we then take the first column of the output and need to call the unique() function? The purpose of the arr.ind argument in the which() function is to return ONLY TRUE values. We take out the first column of the output in order to remove any samples with 0 counts and the unique function to ensure that we don't count any row twice.

We want to filter out any genes (that is the rows) where we have ZERO count data.

```
to.keep inds <- rowSums(meancounts[,1:2] == 0) == 0
head(to.keep inds)
```

```
ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
      TRUE          FALSE          TRUE          TRUE          TRUE
```

```
ENSG000000000938
```

```
      FALSE
```

```
mycounts <- meancounts[to.keep inds,]
nrow(mycounts)
```

```
[1] 21817
```

Q8. Using the up.ind vector above can you determine how many up regulated genes we have at the greater than 2 fc level? There are 250 upregulated genes.

Q9. Using the down.ind vector above can you determine how many down regulated genes we have at the greater than 2 fc level? There are 367 down regulated genes.

A common threshold for calling genes as differentially expressed is a log2 fold-change of +2 or -2.

```
sum(mycounts$log2fc > 2)
```

```
[1] 250
```

```
sum(mycounts$log2fc < -2)
```

```
[1] 367
```

What percent is this?

```
round((sum(mycounts$log2fc >= +2) / nrow(mycounts)) * 100, 2)
```

```
[1] 1.44
```

And down regulated:

```
round((sum(mycounts$log2fc <= -2) / nrow(mycounts)) * 100, 2)
```

```
[1] 2.22
```

Q10. Do you trust these results? Why or why not? I do not trust these results because the data is fold-changed, these values can be large without being statistically significant and we have done nothing to analyze whether these results are statistically significant.

We need some stats to check if the drug induced difference is significant!

Turn to DESeq2

Let's turn to doing this the correct way with the DESeq2 package.

```
library(DESeq2)
```

The main function in the DESeq2 package is called `deseq()`. It wants our count data and our `colData` (metadata) as input in a specific way.

```
dds <- DESeqDataSetFromMatrix(countData = counts,
                               colData = metadata,
                               design = ~dex)
```

converting counts to integer mode

```
Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
design formula are characters, converting to factors
```

Now I can run the DESeq analysis.

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

```
results(dds)
```

```

log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 38694 rows and 6 columns
  baseMean log2FoldChange      lfcSE      stat     pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003  747.1942    -0.3507030  0.168246 -2.084470 0.0371175
ENSG000000000005   0.0000       NA        NA        NA        NA
ENSG000000000419  520.1342    0.2061078  0.101059  2.039475 0.0414026
ENSG000000000457  322.6648    0.0245269  0.145145  0.168982 0.8658106
ENSG000000000460   87.6826    -0.1471420  0.257007 -0.572521 0.5669691
...
...
ENSG00000283115  0.000000       NA        NA        NA        NA
ENSG00000283116  0.000000       NA        NA        NA        NA
ENSG00000283119  0.000000       NA        NA        NA        NA
ENSG00000283120  0.974916    -0.668258   1.69456 -0.394354 0.693319
ENSG00000283123  0.000000       NA        NA        NA        NA
  padj
  <numeric>
ENSG000000000003  0.163035
ENSG000000000005   NA
ENSG000000000419  0.176032
ENSG000000000457  0.961694
ENSG000000000460  0.815849
...
...
ENSG00000283115   NA
ENSG00000283116   NA
ENSG00000283119   NA
ENSG00000283120   NA
ENSG00000283123   NA

```

Now that we have got so far is the log2 fold-change and the adjusted p-value for the significance.

```

res <- results(dds)

head(res)

log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 6 columns
  baseMean log2FoldChange      lfcSE      stat     pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>

```

```

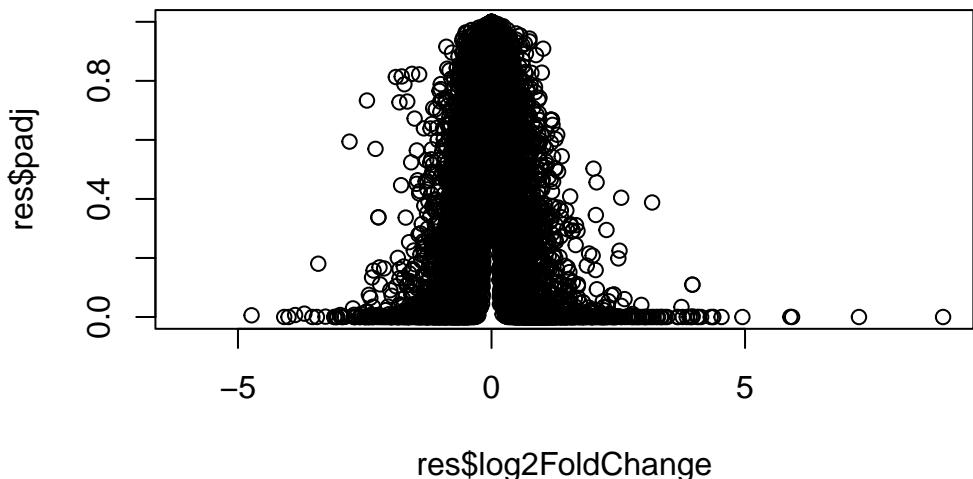
ENSG000000000003 747.194195      -0.3507030  0.168246 -2.084470  0.0371175
ENSG000000000005    0.000000          NA         NA         NA         NA
ENSG000000000419 520.134160      0.2061078  0.101059  2.039475  0.0414026
ENSG000000000457 322.664844      0.0245269  0.145145  0.168982  0.8658106
ENSG000000000460   87.682625     -0.1471420  0.257007 -0.572521  0.5669691
ENSG000000000938   0.319167     -1.7322890  3.493601 -0.495846  0.6200029

      padj
<numeric>
ENSG000000000003  0.163035
ENSG000000000005    NA
ENSG000000000419  0.176032
ENSG000000000457  0.961694
ENSG000000000460  0.815849
ENSG000000000938    NA

```

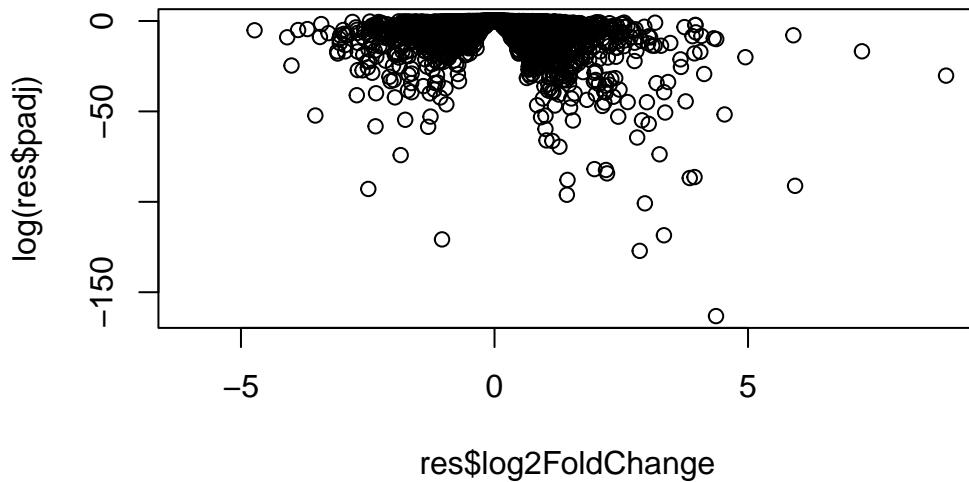
A first plot

```
plot(res$log2FoldChange, res$padj)
```



Well that plot sucked all the interesting P-values are down below zero. I am going to take the log of the p-value

```
plot(res$log2FoldChange, log(res$padj))
```



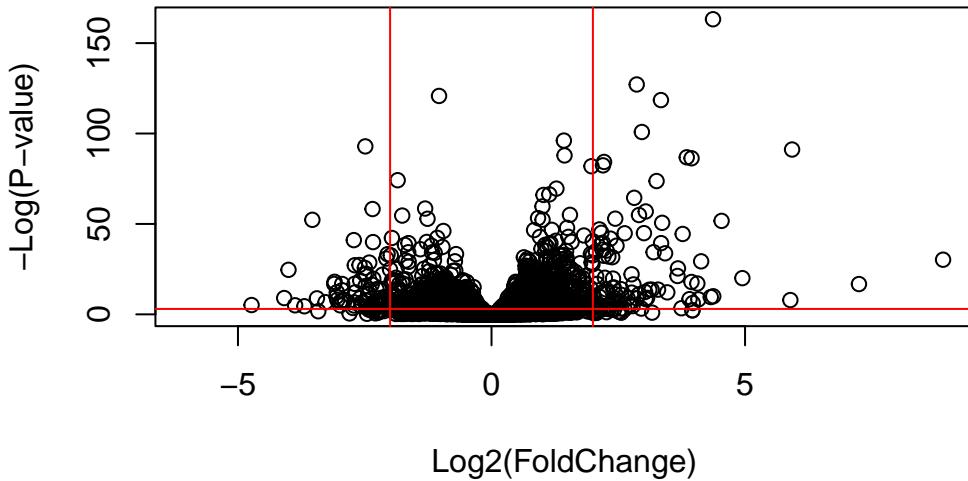
```
log(0.05)
```

```
[1] -2.995732
```

We can flip the y-axis so the plot does not look “upside down”.

```
plot(res$log2FoldChange, -log(res$padj),
     ylab="-Log(P-value)", xlab="Log2(FoldChange)")

abline(v=c(-2,+2), col="red")
abline(h=-log(0.05), col = "red")
```



Add some color to indicate transcripts with large fold change and significant differences between conditions

```
#Setup our custom point color vector
mycols <- rep("gray", nrow(res))
mycols[ abs(res$log2FoldChange) > 2 ] <- "red"

inds <- (res$padj < 0.01) & (abs(res$log2FoldChange) > 2)
mycols[inds] <- "blue"

#Volcano plot with custom colors
plot( res$log2FoldChange, -log(res$padj),
      col=mycols, ylab="-Log(P-value)", xlab="Log2(FoldChange)" )

# Cut off lines
abline(v=c(-2,2), col="gray", lty = 2)
abline(h=-log(0.1), col="gray", lty=2)
```

